

JOINS OUTPUT SCREENSHOTS

```
Activities Terminal Jun 27 16:07
hduser@bmsce-Precision-T1700: ~/Downloads/MapReduceJoin
hduser@bmsce-Precision-T1700:~/Downloads/MapReduceJoin$ hdfs dfs -put ./DeptName.txt /Karthik/dn.txt
put: Parent path is not a directory: /Karthik/Karthik
hduser@bmsce-Precision-T1700:~/Downloads/MapReduceJoin$ hdfs dfs -mkdir /Karthiklast
hduser@bmsce-Precision-T1700:~/Downloads/MapReduceJoin$ hdfs dfs -put ./DeptName.txt /Karthiklast/dn.txt
hduser@bmsce-Precision-T1700:~/Downloads/MapReduceJoin$ hdfs dfs -put ./DeptStrength.txt /Karthiklast/ds.txt
hduser@bmsce-Precision-T1700:~/Downloads/MapReduceJoin$ hadoop jar MapReduceJoin.jar /Karthiklast/ds.txt /Karthiklast/dn.txt /output_karthik_join
22/06/27 16:02:01 INFO Configuration.deprecation: session.id is deprecated. Instead, use dfs.metrics.session-id
22/06/27 16:02:01 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker, sessionId=
22/06/27 16:02:01 INFO jvm.JvmMetrics: Cannot initialize JVM Metrics with processName=JobTracker, sessionId= already initialized
22/06/27 16:02:01 INFO mapred.FileInputFormat: Total input paths to process : 1
22/06/27 16:02:01 INFO mapreduce.JobSubmitter: Cleaning up the staging area file:/app/hadoop/tmp/mapred/staging/hduser154833987/.staging/job_local154833987_0001
Exception in thread "main" org.apache.hadoop.mapred.InvalidInputException: Input path does not exist: hdfs://localhost:54310/user/hduser/dn.txt
    at org.apache.hadoop.mapred.FileInputFormat.singleThreadedListStatus(FileInputFormat.java:285)
    at org.apache.hadoop.mapred.FileInputFormat.listStatus(FileInputFormat.java:228)
    at org.apache.hadoop.mapred.FileInputFormat.getSplits(FileInputFormat.java:313)
    at org.apache.hadoop.mapred.lib.DelegatingInputFormat.getSplits(DelegatingInputFormat.java:107)
    at org.apache.hadoop.mapreduce.JobSubmitter.writeOldSplits(JobSubmitter.java:624)
    at org.apache.hadoop.mapreduce.JobSubmitter.writeSplits(JobSubmitter.java:610)
    at org.apache.hadoop.mapreduce.JobSubmitter.submitJobInternal(JobSubmitter.java:492)
    at org.apache.hadoop.mapreduce.Job$10.run(Job.java:1296)
    at org.apache.hadoop.mapreduce.Job$10.run(Job.java:1293)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:422)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1628)
    at org.apache.hadoop.mapreduce.Job.submit(Job.java:1293)
    at org.apache.hadoop.mapred.JobClient$1.run(JobClient.java:562)
    at org.apache.hadoop.mapred.JobClient$1.run(JobClient.java:557)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:422)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1628)
    at org.apache.hadoop.mapred.JobClient.submitJobInternal(JobClient.java:557)
    at org.apache.hadoop.mapred.JobClient.submitJob(JobClient.java:548)
    at org.apache.hadoop.mapred.JobClient.runJob(JobClient.java:833)
    at MapReduceJoin.JobDriver.run(JobDriver.java:53)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:70)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:84)
    at MapReduceJoin.JobDriver.main(JobDriver.java:68)
    at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
    at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
    at java.lang.reflect.Method.invoke(Method.java:498)
    at org.apache.hadoop.util.RunJar.run(RunJar.java:221)
    at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
hduser@bmsce-Precision-T1700:~/Downloads/MapReduceJoin$ hadoop jar MapReduceJoin.jar /Karthiklast/ds.txt /Karthik/dn.txt /output_karthik_join
22/06/27 16:04:08 INFO Configuration.deprecation: session.id is deprecated. Instead, use dfs.metrics.session-id
22/06/27 16:04:08 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker, sessionId=
22/06/27 16:04:08 INFO jvm.JvmMetrics: Cannot initialize JVM Metrics with processName=JobTracker, sessionId= already initialized
22/06/27 16:04:08 INFO mapreduce.JobSubmitter: Cleaning up the staging area file:/app/hadoop/tmp/mapred/staging/hduser89153233/.staging/job_local89153233_0001
Exception in thread "main" org.apache.hadoop.mapred.InvalidInputException: Input path does not exist: hdfs://localhost:54310/Karthik/dn.txt
    at org.apache.hadoop.mapred.FileInputFormat.singleThreadedListStatus(FileInputFormat.java:285)
    at org.apache.hadoop.mapred.FileInputFormat.listStatus(FileInputFormat.java:228)
    at org.apache.hadoop.mapred.FileInputFormat.getSplits(FileInputFormat.java:313)
    at org.apache.hadoop.mapred.lib.DelegatingInputFormat.getSplits(DelegatingInputFormat.java:107)
    at org.apache.hadoop.mapreduce.JobSubmitter.writeOldSplits(JobSubmitter.java:624)
    at org.apache.hadoop.mapreduce.JobSubmitter.writeSplits(JobSubmitter.java:610)
    at org.apache.hadoop.mapreduce.JobSubmitter.submitJobInternal(JobSubmitter.java:492)
    at org.apache.hadoop.mapreduce.Job$10.run(Job.java:1296)
    at org.apache.hadoop.mapreduce.Job$10.run(Job.java:1293)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:422)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1628)
    at org.apache.hadoop.mapreduce.Job.submit(Job.java:1293)
    at org.apache.hadoop.mapred.JobClient$1.run(JobClient.java:562)
    at org.apache.hadoop.mapred.JobClient$1.run(JobClient.java:557)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:422)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1628)
    at org.apache.hadoop.mapred.JobClient.submitJobInternal(JobClient.java:557)
    at org.apache.hadoop.mapred.JobClient.submitJob(JobClient.java:548)
    at org.apache.hadoop.mapred.JobClient.runJob(JobClient.java:833)
    at MapReduceJoin.JobDriver.run(JobDriver.java:53)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:70)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:84)
    at MapReduceJoin.JobDriver.main(JobDriver.java:68)
    at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
    at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
    at java.lang.reflect.Method.invoke(Method.java:498)
    at org.apache.hadoop.util.RunJar.run(RunJar.java:221)
    at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
```

```
Activities Terminal Jun 27 16:07
hduser@bmsce-Precision-T1700: ~/Downloads/MapReduceJoin
    at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
    at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
    at java.lang.reflect.Method.invoke(Method.java:498)
    at org.apache.hadoop.util.RunJar.run(RunJar.java:221)
    at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
hduser@bmsce-Precision-T1700:~/Downloads/MapReduceJoin$ hadoop jar MapReduceJoin.jar /Karthiklast/ds.txt /Karthik/dn.txt /output_karthik_join
22/06/27 16:05:43 INFO Configuration.deprecation: session.id is deprecated. Instead, use dfs.metrics.session-id
22/06/27 16:05:43 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker, sessionId=
22/06/27 16:05:43 INFO jvm.JvmMetrics: Cannot initialize JVM Metrics with processName=JobTracker, sessionId= already initialized
22/06/27 16:05:43 INFO mapred.FileInputFormat: Total input paths to process : 1
22/06/27 16:05:43 INFO mapreduce.JobSubmitter: number of splits:2
22/06/27 16:05:43 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local1691143096_0001
22/06/27 16:05:43 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
22/06/27 16:05:43 INFO mapred.LocalJobRunner: OutputCommitter set in config null
22/06/27 16:05:43 INFO mapreduce.Job: Running job: job_local1691143096_0001
22/06/27 16:05:43 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapred.FileOutputCommitter
```

```
Activities Terminal Jun 27 16:07
hduser@bmsce-Precision-T1700: ~/Downloads/MapReduceJoin

at org.apache.hadoop.uti.RunJar.run(RunJar.java:221)
at org.apache.hadoop.uti.RunJar.main(RunJar.java:136)
hduser@bmsce-Precision-T1700:~/Downloads/MapReduceJoin$ hadoop jar MapReduceJoin.jar /KarthikIst/ds.txt /KarthikIst/dn.txt /output_karthik_join
22/06/27 16:05:43 INFO Configuration.deprecation: session.id is deprecated. Instead, use dfs.metrics.session-id
22/06/27 16:05:43 INFO jvm.JvmMetrics: Initializing Jvm Metrics with processName=JobTracker, sessionId=
22/06/27 16:05:43 INFO jvm.JvmMetrics: Cannot initialize Jvm Metrics with processName=JobTracker, sessionId= - already initialized
22/06/27 16:05:43 INFO mapred.FileInputFormat: Total input paths to process : 1
22/06/27 16:05:43 INFO mapred.FileInputFormat: Total input paths to process : 1
22/06/27 16:05:43 INFO mapreduce.JobSubmitter: number of splits:2
22/06/27 16:05:43 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local1691143086_0001
22/06/27 16:05:43 INFO mapreduce.Job: The url to track the job: http://localhost:8888/
22/06/27 16:05:43 INFO mapred.LocalJobRunner: OutputCommitter set in config null
22/06/27 16:05:43 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapred.FileOutputCommitter
22/06/27 16:05:43 INFO mapred.LocalJobRunner: Waiting for map tasks
22/06/27 16:05:43 INFO mapred.LocalJobRunner: Starting task: attempt_local1691143086_0001_m_000000_0
22/06/27 16:05:43 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
22/06/27 16:05:43 INFO mapred.MapTask: Processing split: hdfs://localhost:54310/KarthikIst/dn.txt:0+59
22/06/27 16:05:43 INFO mapred.MapTask: numReduceTasks: 1
22/06/27 16:05:43 INFO mapred.MapTask: (EQUATOR) 8 kvt 26214396(104857584)
22/06/27 16:05:43 INFO mapred.MapTask: kvstart = 26214396; length = 6553608
22/06/27 16:05:43 INFO mapred.MapTask: soft limit at 8386088
22/06/27 16:05:43 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857608
22/06/27 16:05:43 INFO mapred.MapTask: kvstart = 26214396; length = 6553608
22/06/27 16:05:43 INFO mapred.MapTask: Map output collector class = org.apache.hadoop.mapred.MapTask$MapOutputBuffer
22/06/27 16:05:43 INFO mapred.LocalJobRunner:
22/06/27 16:05:43 INFO mapred.MapTask: Starting flush of map output
22/06/27 16:05:43 INFO mapred.MapTask: Spilling map output
22/06/27 16:05:43 INFO mapred.MapTask: bufstart = 0; bufend = 63; bufvoid = 104857608
22/06/27 16:05:43 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26214396(104857536); length = 13/6553608
22/06/27 16:05:43 INFO mapred.Task: Task:attempt_local1691143086_0001_m_000000_0 is done. And is in the process of committing
22/06/27 16:05:43 INFO mapred.LocalJobRunner: hdfs://localhost:54310/KarthikIst/dn.txt:0+59
22/06/27 16:05:43 INFO mapred.Task: Task:attempt_local1691143086_0001_m_000000_0 done.
22/06/27 16:05:43 INFO mapred.LocalJobRunner: Finishing task: attempt_local1691143086_0001_m_000000_0
22/06/27 16:05:43 INFO mapred.LocalJobRunner: Starting task: attempt_local1691143086_0001_m_000001_0
22/06/27 16:05:43 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
22/06/27 16:05:43 INFO mapred.MapTask: Processing split: hdfs://localhost:54310/KarthikIst/ds.txt:0+58
22/06/27 16:05:43 INFO mapred.MapTask: (EQUATOR) 8 kvt 26214396(104857584)
22/06/27 16:05:43 INFO mapred.MapTask: kvstart = 26214396; length = 6553608
22/06/27 16:05:43 INFO mapred.MapTask: soft limit at 8386088
22/06/27 16:05:43 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857608
22/06/27 16:05:43 INFO mapred.MapTask: kvstart = 26214396; length = 6553608
22/06/27 16:05:43 INFO mapred.MapTask: Map output collector class = org.apache.hadoop.mapred.MapTask$MapOutputBuffer
22/06/27 16:05:43 INFO mapred.LocalJobRunner:
22/06/27 16:05:43 INFO mapred.MapTask: Starting flush of map output
22/06/27 16:05:43 INFO mapred.MapTask: Spilling map output
22/06/27 16:05:43 INFO mapred.MapTask: bufstart = 0; bufend = 54; bufvoid = 104857608
22/06/27 16:05:43 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26214384(104857536); length = 13/6553608
22/06/27 16:05:43 INFO mapred.Task: Task:attempt_local1691143086_0001_m_000001_0 is done. And is in the process of committing
22/06/27 16:05:43 INFO mapred.LocalJobRunner: hdfs://localhost:54310/KarthikIst/ds.txt:0+58
22/06/27 16:05:43 INFO mapred.Task: Task:attempt_local1691143086_0001_m_000001_0 done.
22/06/27 16:05:43 INFO mapred.LocalJobRunner: Finishing task: attempt_local1691143086_0001_m_000001_0
22/06/27 16:05:43 INFO mapred.LocalJobRunner: map task executor complete.
22/06/27 16:05:44 INFO mapred.LocalJobRunner: map task executor complete.
```

```
Activities Terminal Jun 27 16:07
hduser@bmsce-Precision-T1700: ~/Downloads/MapReduceJoin

22/06/27 16:05:43 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
22/06/27 16:05:43 INFO mapred.MapTask: Processing split: hdfs://localhost:54310/KarthikIst/ds.txt:0+58
22/06/27 16:05:43 INFO mapred.MapTask: numReduceTasks: 1
22/06/27 16:05:43 INFO mapred.MapTask: (EQUATOR) 8 kvt 26214396(104857584)
22/06/27 16:05:43 INFO mapred.MapTask: kvstart = 26214396; length = 6553608
22/06/27 16:05:43 INFO mapred.MapTask: soft limit at 8386088
22/06/27 16:05:43 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857608
22/06/27 16:05:43 INFO mapred.MapTask: kvstart = 26214396; length = 6553608
22/06/27 16:05:43 INFO mapred.MapTask: Map output collector class = org.apache.hadoop.mapred.MapTask$MapOutputBuffer
22/06/27 16:05:43 INFO mapred.LocalJobRunner:
22/06/27 16:05:43 INFO mapred.MapTask: Starting flush of map output
22/06/27 16:05:43 INFO mapred.MapTask: Spilling map output
22/06/27 16:05:43 INFO mapred.MapTask: bufstart = 0; bufend = 54; bufvoid = 104857608
22/06/27 16:05:43 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26214384(104857536); length = 13/6553608
22/06/27 16:05:43 INFO mapred.Task: Task:attempt_local1691143086_0001_m_000001_0 is done. And is in the process of committing
22/06/27 16:05:43 INFO mapred.LocalJobRunner: hdfs://localhost:54310/KarthikIst/ds.txt:0+58
22/06/27 16:05:43 INFO mapred.Task: Task:attempt_local1691143086_0001_m_000001_0 done.
22/06/27 16:05:43 INFO mapred.LocalJobRunner: Finishing task: attempt_local1691143086_0001_m_000001_0
22/06/27 16:05:43 INFO mapred.LocalJobRunner: map task executor complete.
22/06/27 16:05:44 INFO mapred.LocalJobRunner: Waiting for reduce tasks
22/06/27 16:05:44 INFO mapred.LocalJobRunner: Starting task: attempt_local1691143086_0001_r_000000_0
22/06/27 16:05:44 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
22/06/27 16:05:44 INFO mapred.ReduceTask: Using ShuffleConsumerPlugin: org.apache.hadoop.mapreduce.task.reduce.Shuffle7d9d9d08
22/06/27 16:05:44 INFO reduce.MergeManagerImpl: MergeManager: memoryLimit=33438464, maxSingleShuffleLimit=33584616, mergeThreshold=32863392, sortFactor=16, nextNonMergeOutputsThreshold=10
22/06/27 16:05:44 INFO reduce.EventFetcher: attempt_local1691143086_0001_r_000000_0 Thread started: EventFetcher for fetching Map Completion Events
22/06/27 16:05:44 INFO reduce.LocalFetcher: about to shuffle output of map attempt_local1691143086_0001_m_000001_0 decom: 64 len: 68 to MEMORY
22/06/27 16:05:44 INFO reduce.InMemoryMapOutput: Read 64 bytes from map-output for attempt_local1691143086_0001_m_000001_0
22/06/27 16:05:44 INFO reduce.MergeManagerImpl: closeInMemoryFile -> map-output of size: 64, inMemoryMapOutputs.size() -> 1, currInMemory -> 8, usedMemory -> 64
22/06/27 16:05:44 INFO reduce.LocalFetcher: LocalFetcher: about to shuffle output of map attempt_local1691143086_0001_m_000001_0 decom: 73 len: 77 to MEMORY
22/06/27 16:05:44 INFO reduce.InMemoryMapOutput: Read 73 bytes from map-output for attempt_local1691143086_0001_m_000001_0
22/06/27 16:05:44 INFO reduce.MergeManagerImpl: closeInMemoryFile -> map-output of size: 73, inMemoryMapOutputs.size() -> 2, currInMemory -> 84, usedMemory -> 137
22/06/27 16:05:44 INFO reduce.EventFetcher: EventFetcher is interrupted.. Returning
22/06/27 16:05:44 INFO mapred.LocalJobRunner: 2 / 2 copied.
22/06/27 16:05:44 INFO reduce.MergeManagerImpl: finalMerge called with 2 in memory map-outputs and 0 on-disk map-outputs
22/06/27 16:05:44 INFO mapred.Merger: Merging 2 sorted segments
22/06/27 16:05:44 INFO mapred.Merger: Down to the last merge pass, with 2 segments left of total size: 121 bytes
22/06/27 16:05:44 INFO reduce.MergeManagerImpl: Merged 2 segments, 137 bytes to disk to satisfy reduce memory limit
22/06/27 16:05:44 INFO reduce.MergeManagerImpl: Merging 1 files, 139 bytes from disk
22/06/27 16:05:44 INFO reduce.MergeManagerImpl: Merging 0 segments, 0 bytes from memory into reduce
22/06/27 16:05:44 INFO mapred.Merger: Merging 1 sorted segments
22/06/27 16:05:44 INFO mapred.Merger: Down to the last merge pass, with 1 segment left of total size: 127 bytes
22/06/27 16:05:44 INFO mapred.LocalJobRunner: 2 / 2 copied.
22/06/27 16:05:44 INFO mapred.Task: Task:attempt_local1691143086_0001_r_000000_0 is allowed to commit now
22/06/27 16:05:44 INFO mapred.Task: Task attempt_local1691143086_0001_r_000000_0 to hdfs://localhost:54310/output_karthik_join/_temporary/0/task_local1691143086_0001_r_000000_0
22/06/27 16:05:44 INFO output.FileOutputCommitter: Saved output of task 'attempt_local1691143086_0001_r_000000_0' to hdfs://localhost:54310/output_karthik_join/_temporary/0/task_local1691143086_0001_r_000000_0
22/06/27 16:05:44 INFO mapred.LocalJobRunner: reduce > reduce
22/06/27 16:05:44 INFO mapred.Task: Task 'attempt_local1691143086_0001_r_000000_0' done.
22/06/27 16:05:44 INFO mapred.LocalJobRunner: Finishing task: attempt_local1691143086_0001_r_000000_0
22/06/27 16:05:44 INFO mapred.LocalJobRunner: reduce task executor complete.
22/06/27 16:05:44 INFO mapreduce.Job: Job job_local1691143086_0001 running in uber mode : false
22/06/27 16:05:44 INFO mapreduce.Job: map 100% reduce 100%
22/06/27 16:05:44 INFO mapreduce.Job: Job job_local1691143086_0001 completed successfully
```

```
Activities Terminal Jun 27 16:07
hduser@bmsce-Precision-T1700: ~/Downloads/MapReduceJoin

22/06/27 16:05:44 INFO mapred.LocalJobRunner: reduce > reduce
22/06/27 16:05:44 INFO mapred.Task: Task 'attempt_local1691143806_0001_r_000000_0' done.
22/06/27 16:05:44 INFO mapred.LocalJobRunner: finishing task: attempt_local1691143806_0001_r_000000_0
22/06/27 16:05:44 INFO mapred.LocalJobRunner: reduce task executor complete.
22/06/27 16:05:44 INFO mapreduce.Job: Job job_local1691143806_0001 running in uber mode : false
22/06/27 16:05:44 INFO mapreduce.Job: map 100% reduce 100%
22/06/27 16:05:44 INFO mapreduce.Job: Job job_local1691143806_0001 completed successfully
22/06/27 16:05:44 INFO mapreduce.Job: Counters: 38

File System Counters
  FILE: Number of bytes read=26288
  FILE: Number of bytes written=782595
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=277
  HDFS: Number of bytes written=85
  HDFS: Number of read operations=28
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=5

Map-Reduce Framework
  Map input records=8
  Map output records=8
  Map output bytes=117
  Map output materialized bytes=145
  Input split bytes=495
  Combine input records=0
  Combine output records=0
  Reduce input groups=1
  Reduce shuffle bytes=145
  Reduce input records=8
  Reduce output records=4
  Spilled Records=16
  Shuffled Maps=2
  Failed Shuffles=0
  Merged Map outputs=2
  GC time elapsed (ms)=1
  CPU time spent (ms)=8
  Physical memory (bytes) snapshot=0
  Virtual memory (bytes) snapshot=0
  Total committed heap usage (bytes)=923271168

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=0
  File Output Format Counters
    Bytes Written=85

hduser@bmsce-Precision-T1700:~/Downloads/MapReduceJoin$ hadoop fs -ls /
found 46 items
drwxr-xr-x  - hduser supergroup      0 2022-06-27 11:09 /Akshatha
-rw-r--r--  1 hduser supergroup    888190 2022-06-27 15:10 /Karthik
drwxr-xr-x  - hduser supergroup      0 2022-06-27 15:58 /Karthiklast
```