

TEMPERATURE PROGRAM OUTPUT SCREENSHOTS

```
Activities Terminal Jun 27 15:50
hduser@bmisce-Precision-T1700:~$
hduser@bmisce-Precision-T1700:~$ hadoop hduser
Error: Could not find or load main class hduser
hduser@bmisce-Precision-T1700:~$ hadoop
Usage: hadoop [-config confdir] COMMAND
  where COMMAND is one of:
    fs                run a generic filesystem user client
    version            print the version
    jar <jar>         run a jar file
    checknative [-a|-s] check native hadoop and compression libraries availability
    distcp <srcurl> <dsturl> copy file or directories recursively
    archive <archiveName> <name> <p> <parent path> <srcs> <dsts> create a hadoop archive
    classpath          prints the class path needed to get the
                      interact with credential providers
    credential         Hadoop jar and the required libraries
    daemonlog          get/set the log level for each daemon
    trace             view and modify Hadoop tracing settings
    or
    CLASSNAME         run the class named CLASSNAME

Most commands print help when invoked w/o parameters.
hduser@bmisce-Precision-T1700:~$ start-all.sh
This script is deprecated. Instead use start-dfs.sh and start-yarn.sh
Starting namenodes on [localhost]
hduser@localhost's password:
hduser@localhost's password:
localhost: starting datanode, logging to /usr/local/hadoop/logs/hadoop-hduser-datanode-bmisce-Precision-T1700.out
Starting secondary namenodes [0.0.0.0]
hduser@0.0.0.0's password:
0.0.0.0: starting secondarynamenode, logging to /usr/local/hadoop/logs/hadoop-hduser-secondarynamenode-bmisce-Precision-T1700.out
starting yarn daemons
starting resourcemanager, logging to /usr/local/hadoop/logs/yarn-hduser-resourcemanager-bmisce-Precision-T1700.out
hduser@localhost's password:
localhost: starting nodemanager, logging to /usr/local/hadoop/logs/yarn-hduser-nodemanager-bmisce-Precision-T1700.out
hduser@bmisce-Precision-T1700:~$ jps
6176 ResourceManager
6817 SecondaryNameNode
5623 NameNode
6633 Jps
5905 DataNode
4382 org.eclipse.equinox.launcher_1.5.000.v20191014-2022.jar
6511 NodeManager
hduser@bmisce-Precision-T1700:~$ hdfs dfs -mkdir /input_dir
hduser@bmisce-Precision-T1700:~$ hdfs dfs -ls /
Found 36 items
drwxr-xr-x - hduser supergroup 0 2022-06-27 13:09 /akshatha
drwxr-xr-x - hduser supergroup 0 2022-06-26 12:21 /jill
drwxr-xr-x - hduser supergroup 0 2022-06-24 15:26 /jabc
drwxr-xr-x - hduser supergroup 0 2022-06-03 12:18 /jabcd
drwxr-xr-x - hduser supergroup 0 2022-05-31 18:12 /janother
-rw-r--r-- 1 hduser supergroup 888976 2022-06-22 14:41 /avgtemp_outputdir
drwxr-xr-x - hduser supergroup 0 2022-06-22 15:52 /avgtempdir
drwxr-xr-x - hduser supergroup 0 2022-06-01 15:22 /dest
drwxr-xr-x - hduser supergroup 0 2022-06-22 15:52 /final_output
drwxr-xr-x - hduser supergroup 0 2022-06-22 15:27 /final_output1
drwxr-xr-x - hduser supergroup 0 2022-06-22 15:46 /final_output2
```

```
Activities Terminal Jun 27 15:50
hduser@bmisce-Precision-T1700:~$
drwxr-xr-x - hduser supergroup 0 2022-06-25 18:46 /joel
drwxr-xr-x - hduser supergroup 0 2022-06-20 15:16 /kiran
drwxr-xr-x - hduser supergroup 0 2022-06-24 15:26 /lab
drwxr-xr-x - hduser supergroup 0 2022-06-27 12:13 /new
drwxr-xr-x - hduser supergroup 0 2019-10-24 10:58 /nlldh
drwxr-xr-x - hduser supergroup 0 2022-06-23 19:15 /output
drwxr-xr-x - hduser supergroup 0 2022-06-27 12:52 /output_111
drwxr-xr-x - hduser supergroup 0 2022-06-25 18:30 /output_199
drwxr-xr-x - hduser supergroup 0 2022-06-26 15:18 /output_kiran
drwxr-xr-x - hduser supergroup 0 2022-06-15 14:47 /output_nitth
drwxr-xr-x - hduser supergroup 0 2022-06-15 14:48 /output_nitth1
drwxr-xr-x - hduser supergroup 0 2022-06-15 15:06 /output_nitth3
drwxr-xr-x - hduser supergroup 0 2022-06-15 15:18 /output_nitth4
drwxr-xr-x - hduser supergroup 0 2022-06-22 15:22 /output_nitthn
drwxr-xr-x - hduser supergroup 0 2022-05-31 09:43 /poolja
-rw-r--r-- 1 hduser supergroup 12 2022-05-31 10:25 /pooljah
drwxr-xr-x - hduser supergroup 0 2022-06-15 14:46 /qqs
drwxr-xr-x - hduser supergroup 0 2022-06-01 14:09 /temp
drwxr-xr-x - hduser supergroup 0 2019-08-01 10:19 /tmp
drwxr-xr-x - hduser supergroup 0 2022-06-01 09:46 /tusher
drwxr-xr-x - hduser supergroup 0 2022-06-20 12:26 /user
hduser@bmisce-Precision-T1700:~$ hdfs sfs -ls /input_dir
Error: Could not find or load main class sfs
hduser@bmisce-Precision-T1700:~$ hdfs dfs -ls /input_dir
hduser@bmisce-Precision-T1700:~$ hdfs dfs -copyFromLocal C:\input.txt /input_dir
-copiedFromLocal: Can not create a Path from a null string
Usage: hadoop fs [generic options] -copyFromLocal [-f] [-p] [-l] <localsrc> ... <dst>
hduser@bmisce-Precision-T1700:~$ hdfs dfs -ls /input_dir
Found 1 items
-rw-r--r-- 1 hduser supergroup 9 2022-06-27 14:54 /input_dir/local.txt
hduser@bmisce-Precision-T1700:~$ hdfs dfs -cat /input_dir/local.txt
cat: '/input_dir/local.txt': No such file or directory
hduser@bmisce-Precision-T1700:~$ hdfs dfs -cat /input_dir/local.txt
Hi Hello
hduser@bmisce-Precision-T1700:~$ hadoop jar /hduser/Desktop/Karthik.jar TopN /input_dir/local.txt /output_karthik
Not a valid JAR: /hduser/Desktop/Karthik.jar
hduser@bmisce-Precision-T1700:~$ hadoop jar /home/hduser/Desktop/Karthik.jar TopN /input_dir/local.txt /output_karthik
Not a valid JAR: /home/hduser/Desktop/Karthik.jar
hduser@bmisce-Precision-T1700:~$ hadoop jar /home/hduser/Desktop/Karthik.jar TopN /input_dir/local.txt /output_karthik
Exception in thread "main" java.lang.ClassNotFoundException: TopN
at java.net.URLClassLoader.findClass(URLClassLoader.java:382)
at java.lang.ClassLoader.loadClass(ClassLoader.java:424)
at java.lang.ClassLoader.loadClass(ClassLoader.java:357)
at java.lang.Class.forName0(Native Method)
at java.lang.Class.forName(Class.java:348)
at org.apache.hadoop.util.RunJar.run(RunJar.java:214)
at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
hduser@bmisce-Precision-T1700:~$ hadoop jar /home/hduser/Desktop/Karthik.jar Karthik.TopN /input_dir/local.txt /output_karthik
22/06/27 15:55:49 INFO Configuration.deprecation: session-id is deprecated. Instead, use dfs.metrics.session-id
22/06/27 15:55:49 INFO jvm.JVMMetrics: Initializing JVM Metrics with processName=JobTracker, sessionId=
22/06/27 15:55:49 INFO InputFileInputFormat: Total input paths to process : 1
22/06/27 15:55:49 INFO mapreduce.JobSubmitter: number of splits:1
22/06/27 15:55:49 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local91709511_0001
22/06/27 15:55:49 INFO mapreduce.Job: The url to track the job: http://localhost:8088/
22/06/27 15:55:49 INFO mapreduce.Job: Running job: job_local91709511_0001
```

```
at java.lang.ClassLoader.loadClass(ClassLoader.java:424)
at java.lang.ClassLoader.loadClass(ClassLoader.java:357)
at java.lang.Class.forName0(Native Method)
at java.lang.Class.forName(Class.java:348)
at org.apache.hadoop.util.RunJar.run(RunJar.java:214)
at org.apache.hadoop.util.RunJar.main(RunJar.java:130)
hduser@bmsce-Precision-T1700:~$ hadoop jar /home/hduser/Desktop/Karthik.jar Karthik.topM /input_dir/local.txt /output_karthik
22/06/27 15:05:49 INFO Configuration.deprecation: session.id is deprecated. Instead, use dfs.metrics.session-id
22/06/27 15:05:49 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker, sessionId=
22/06/27 15:05:49 INFO InputFileInputFormat: Total input paths to process : 1
22/06/27 15:05:49 INFO mapreduce.JobSubmitter: number of splits=1
22/06/27 15:05:49 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local01709511_0001
22/06/27 15:05:49 INFO mapreduce.Job: The url to track the job: http://localhost:8888/
22/06/27 15:05:49 INFO mapreduce.Job: RunningJob: job_local01709511_0001
22/06/27 15:05:49 INFO mapreduce.LocalJobRunner: OutputCommitter set in config null
22/06/27 15:05:49 INFO mapreduce.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
22/06/27 15:05:49 INFO mapreduce.LocalJobRunner: Waiting for map tasks
22/06/27 15:05:49 INFO mapreduce.LocalJobRunner: Starting task: attempt_local01709511_0001_r_000000_0
22/06/27 15:05:49 INFO mapreduce.Task: Using ResourceCalculatorProcessTree : [ ]
22/06/27 15:05:49 INFO mapreduce.MapTask: Processing split: hdfs://localhost:54310/input_dir/local.txt:0+9
22/06/27 15:05:50 INFO mapreduce.MapTask: (EQUIOR) 0 kv[ 26214396(104857584)
22/06/27 15:05:50 INFO mapreduce.MapTask: mapreduce.task.io.sort.mb: 100
22/06/27 15:05:50 INFO mapreduce.MapTask: sort limit at 83886080
22/06/27 15:05:50 INFO mapreduce.MapTask: burststart = 0; bufvoid = 104857600
22/06/27 15:05:50 INFO mapreduce.MapTask: kvstart = 26214396; length = 6553600
22/06/27 15:05:50 INFO mapreduce.MapTask: Map output collector class = org.apache.hadoop.mapreduce.MapTask$MapOutputBuffer
22/06/27 15:05:50 INFO mapreduce.LocalJobRunner:
22/06/27 15:05:50 INFO mapreduce.MapTask: Starting flush of map output
22/06/27 15:05:50 INFO mapreduce.MapTask: Spilling map output
22/06/27 15:05:50 INFO mapreduce.MapTask: burststart = 0; bufvoid = 17; bufvoid = 104857600
22/06/27 15:05:50 INFO mapreduce.MapTask: kvstart = 26214396(104857584); kvend = 26214392(104857568); length = 5/6553600
22/06/27 15:05:50 INFO mapreduce.MapTask: Finished spill 0
22/06/27 15:05:50 INFO mapreduce.Task: Task attempt_local01709511_0001_r_000000_0 is done. And is in the process of committing
22/06/27 15:05:50 INFO mapreduce.LocalJobRunner: map
22/06/27 15:05:50 INFO mapreduce.Task: Task attempt_local01709511_0001_r_000000_0 done.
22/06/27 15:05:50 INFO mapreduce.LocalJobRunner: Finishing task: attempt_local01709511_0001_r_000000_0
22/06/27 15:05:50 INFO mapreduce.LocalJobRunner: map task executor complete.
22/06/27 15:05:50 INFO mapreduce.LocalJobRunner: Waiting for reduce tasks
22/06/27 15:05:50 INFO mapreduce.LocalJobRunner: Starting task: attempt_local01709511_0001_r_000000_0
22/06/27 15:05:50 INFO mapreduce.Task: Using ResourceCalculatorProcessTree : [ ]
22/06/27 15:05:50 INFO mapreduce.ReduceTask: Using ShuffleCommitterImpl: org.apache.hadoop.mapreduce.Task.reduce.ShuffleCommitter
22/06/27 15:05:50 INFO reduce.MergeManagerImpl: MergeManager: memoryLimit=134338404, maxSingleShuffleLimit=83886016, mergeThreshold=220003392, ioSortFactor=10, mergeOutputThreshold=10
22/06/27 15:05:50 INFO reduce.EventFetcher: attempt_local01709511_0001_r_000000_0 thread started: EventFetcher for fetching Map completion events
22/06/27 15:05:50 INFO reduce.EventFetcher: LocalFetcher: about to shuffle output of map attempt_local01709511_0001_r_000000_0 decom: 23 len: 27 to MEMORY
22/06/27 15:05:50 INFO reduce.InMemoryMapOutput: Read 23 bytes from map-output for attempt_local01709511_0001_r_000000_0
22/06/27 15:05:50 INFO reduce.MergeManagerImpl: closeInMemoryFile -> map-output of size: 23, inMemoryMapOutputs.size() -> 1, commitMemory -> 0, usedMemory -> 23
22/06/27 15:05:50 INFO reduce.EventFetcher: EventFetcher is interrupted.. Returning
22/06/27 15:05:50 INFO mapreduce.LocalJobRunner: 1 / 1 copied.
22/06/27 15:05:50 INFO reduce.MergeManagerImpl: finalMerge called with 1 in-memory map-outputs and 0 on-disk map-outputs
22/06/27 15:05:50 INFO mapreduce.Reducer: Merging 1 sorted segments
22/06/27 15:05:50 INFO mapreduce.Reducer: Down to the last merge-pass, with 1 segments left of total size: 15 bytes
22/06/27 15:05:50 INFO reduce.MergeManagerImpl: Merging 1 segments, 22 bytes to disk to satisfy reduce memory limit
22/06/27 15:05:50 INFO reduce.MergeManagerImpl: Merging 1 files, 27 bytes from disk
22/06/27 15:05:50 INFO reduce.MergeManagerImpl: Merging 0 segments, 0 bytes from memory into reduce
22/06/27 15:05:50 INFO mapreduce.Reducer: Merging 1 sorted segments
22/06/27 15:05:50 INFO mapreduce.Reducer: Down to the last merge-pass, with 1 segments left of total size: 15 bytes
```

```
22/06/27 15:05:50 INFO mapreduce.Job: Job job_local01709511_0001 running in uber mode : false
22/06/27 15:05:50 INFO mapreduce.Job: map 100% reduce 100%
22/06/27 15:05:50 INFO mapreduce.Job: Job job_local01709511_0001 completed successfully
22/06/27 15:05:50 INFO mapreduce.Job: Counters: 38
File System Counters
  FILE: Number of bytes read=20370
  FILE: Number of bytes written=521101
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=16
  HDFS: Number of bytes written=13
  HDFS: Number of read operations=13
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=4
Map-Reduce Framework
  Map input records=1
  Map output records=2
  Map output materialized bytes=27
  Input split bytes=107
  Combine input records=0
  Combine output records=0
  Reduce input groups=2
  Reduce shuffle bytes=27
  Reduce input records=2
  Reduce output records=2
  Spilled Records=4
  Shuffled Maps=1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=41
  CPU time spent (ms)=0
  Physical memory (bytes) snapshot=0
  Virtual memory (bytes) snapshot=0
  Total committed heap usage (bytes)=571908788
Shuffle
  Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=9
File Output Format Counters
  Bytes Written=13
hduser@bmsce-Precision-T1700:~$ hadoop fs -ls /
Found 37 items
drwxr-xr-x - hduser supergroup 0 2022-06-27 13:09 /akshatha
drwxr-xr-x - hduser supergroup 0 2022-06-06 12:21 /llll
drwxr-xr-x - hduser supergroup 0 2022-06-06 14:50 /abc
drwxr-xr-x - hduser supergroup 0 2022-06-03 12:18 /abcd
drwxr-xr-x - hduser supergroup 0 2022-05-31 10:12 /another
-rw-r--r-- 1 hduser supergroup 888976 2022-06-22 14:11 /avgtemp_outputdir
drwxr-xr-x - hduser supergroup 0 2022-06-22 14:52 /avgtempdir
```

```
Activities Terminal Jun 27 15:50
hduser@bmisce-Precision-T1700: ~
Found 37 items
drwxr-xr-x - hduser supergroup 0 2022-06-27 13:09 /kshatha
drwxr-xr-x - hduser supergroup 0 2022-06-06 12:21 /LL
drwxr-xr-x - hduser supergroup 0 2022-06-06 14:56 /abc
drwxr-xr-x - hduser supergroup 0 2022-06-03 12:18 /abcd
drwxr-xr-x - hduser supergroup 0 2022-05-21 13:12 /another
-rw-r--r-- - hduser supergroup 888978 2022-06-22 14:41 /avgtemp_outputdir
drwxr-xr-x - hduser supergroup 0 2022-06-22 14:52 /avgtempdir
drwxr-xr-x - hduser supergroup 0 2022-06-01 15:22 /dest
drwxr-xr-x - hduser supergroup 0 2022-06-22 14:52 /final_output
drwxr-xr-x - hduser supergroup 0 2022-06-22 15:27 /final_output1
drwxr-xr-x - hduser supergroup 0 2022-06-22 15:46 /final_output2
drwxr-xr-x - hduser supergroup 0 2022-06-01 09:32 /hadoop
drwxr-xr-x - hduser supergroup 0 2022-06-20 12:25 /hadoop1
drwxr-xr-x - hduser supergroup 0 2022-06-27 14:54 /input_dir
drwxr-xr-x - hduser supergroup 0 2022-06-22 15:45 /inputnithin
drwxr-xr-x - hduser supergroup 0 2022-06-25 15:46 /joel
drwxr-xr-x - hduser supergroup 0 2022-06-20 15:10 /kiran
drwxr-xr-x - hduser supergroup 0 2022-06-24 15:29 /lab
drwxr-xr-x - hduser supergroup 0 2022-06-27 12:13 /new
drwxr-xr-x - hduser supergroup 0 2019-10-24 13:58 /ntdu
drwxr-xr-x - hduser supergroup 0 2022-06-25 18:15 /output
drwxr-xr-x - hduser supergroup 0 2022-06-27 12:52 /output_111
drwxr-xr-x - hduser supergroup 0 2022-06-25 18:30 /output_199
drwxr-xr-x - hduser supergroup 0 2022-06-27 15:05 /output_karthik
drwxr-xr-x - hduser supergroup 0 2022-06-20 15:18 /output_kiran
drwxr-xr-x - hduser supergroup 0 2022-06-15 14:47 /output_nitin
drwxr-xr-x - hduser supergroup 0 2022-06-15 14:48 /output_nitin1
drwxr-xr-x - hduser supergroup 0 2022-06-15 15:06 /output_nitin3
drwxr-xr-x - hduser supergroup 0 2022-06-15 15:18 /output_nitin4
drwxr-xr-x - hduser supergroup 0 2022-06-22 15:22 /outputnithin
drwxr-xr-x - hduser supergroup 0 2022-05-31 09:43 /poola
-rw-r--r-- - hduser supergroup 12 2022-05-31 18:25 /pooljah
drwxr-xr-x - hduser supergroup 0 2022-06-15 14:40 /rgs
drwxr-xr-x - hduser supergroup 0 2022-06-01 14:40 /temp
drwxr-xr-x - hduser supergroup 0 2019-08-01 16:19 /tmp
drwxr-xr-x - hduser supergroup 0 2022-06-01 09:46 /tushar
drwxr-xr-x - hduser supergroup 0 2022-06-20 12:28 /user
hduser@bmisce-Precision-T1700: ~$ hadoop fs -ls /output_karthik
Found 2 items
-rw-r--r-- 1 hduser supergroup 0 2022-06-27 15:05 /output_karthik/ SUCCESS
-rw-r--r-- 1 hduser supergroup 13 2022-06-27 15:05 /output_karthik/part-r-00000
hduser@bmisce-Precision-T1700: ~$ hadoop fs -cat /output_karthik/part-r-00000
hl
1
hello
1
hduser@bmisce-Precision-T1700: ~$ hadoop fs -copyFromLocal /home/hduser/Downloads/1981.txt /Karthik
hduser@bmisce-Precision-T1700: ~$ hadoop fs -copyFromLocal /home/hduser/Desktop/Karthiktemp.jar temp.AverageDriver /Karthik/1981.txt
hduser@bmisce-Precision-T1700: ~$ hadoop jar /home/hduser/Desktop/Karthiktemp.jar temp.AverageDriver /Karthik/1981.txt /output_karthiktemp
Please Enter the input and output parameters
22/06/27 15:20:11 INFO Configuration.deprecation: session.id is deprecated. Instead, use dfs.metrics.session-id
22/06/27 15:20:11 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker, sessionId=
22/06/27 15:20:11 WARN MapReduce.JobSubmitter: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
22/06/27 15:20:11 INFO mapreduce.JobSubmitter: Cleaning up the staging area file:/app/hadoop/nu/mapred/staging/hduser111999/26/.staging/job_local111999/36_0001
Exception in thread "main" org.apache.hadoop.mapreduce.lib.input.InvalidInputException: Input path does not exist: hdfs://localhost:54310/Karthik/1981.txt
at org.apache.hadoop.mapreduce.lib.input.FileInputFormat.singleThreadedListStatus(FileInputFormat.java:321)
```

```
Activities Terminal Jun 27 15:51
hduser@bmisce-Precision-T1700: ~
-rw-r--r-- 1 hduser supergroup 12 2022-05-31 18:25 /pooljah
drwxr-xr-x - hduser supergroup 0 2022-06-15 14:40 /rgs
drwxr-xr-x - hduser supergroup 0 2022-06-01 14:40 /temp
drwxr-xr-x - hduser supergroup 0 2019-08-01 16:19 /tmp
drwxr-xr-x - hduser supergroup 0 2022-06-01 09:46 /tushar
drwxr-xr-x - hduser supergroup 0 2022-06-20 12:28 /user
hduser@bmisce-Precision-T1700: ~$ hadoop fs -ls /Karthiknew
Found 1 items
-rw-r--r-- 1 hduser supergroup 888190 2022-06-27 15:23 /Karthiknew/1901.txt
hduser@bmisce-Precision-T1700: ~$ hadoop fs -ls /output_karthiktemp
Found 3 items
-rw-r--r-- 1 hduser supergroup 0 2022-06-27 15:23 /output_karthiktemp/ SUCCESS
-rw-r--r-- 1 hduser supergroup 8 2022-06-27 15:23 /output_karthiktemp/part-r-00000
hduser@bmisce-Precision-T1700: ~$ hadoop fs -ls /output_karthiktemp
Found 2 items
-rw-r--r-- 1 hduser supergroup 0 2022-06-27 15:23 /output_karthiktemp/ SUCCESS
-rw-r--r-- 1 hduser supergroup 0 2022-06-27 15:23 /output_karthiktemp/part-r-00000
hduser@bmisce-Precision-T1700: ~$ hadoop fs -ls /output_karthiktemp/part-r-00000
-rw-r--r-- 1 hduser supergroup 8 2022-06-27 15:23 /output_karthiktemp/part-r-00000
hduser@bmisce-Precision-T1700: ~$ hadoop -cat /output_karthiktemp/part-r-00000
Error: No command named '-cat' was found. Perhaps you meant 'hadoop cat'
hduser@bmisce-Precision-T1700: ~$ hdfs dfs -cat /output_karthiktemp/part-r-00000
1981 46
hduser@bmisce-Precision-T1700: ~$ hadoop fs -copyFromLocal /home/hduser/Downloads/1982.txt /Karthiktemp
copyFromLocal: /Karthiktemp: File exists
hduser@bmisce-Precision-T1700: ~$ hadoop fs -copyFromLocal /home/hduser/Downloads/1982.txt /Karthiktemp2
copyFromLocal: /home/hduser/Downloads/1982.txt: No such file or directory
hduser@bmisce-Precision-T1700: ~$ hadoop fs -copyFromLocal /home/hduser/Downloads/1982.txt /Karthiktemp
copyFromLocal: /Karthiktemp: File exists
hduser@bmisce-Precision-T1700: ~$ hadoop fs -copyFromLocal /home/hduser/Downloads/1982.txt /Karthiktemp2
hduser@bmisce-Precision-T1700: ~$ hadoop jar /home/hduser/Desktop/Karthiktemp.jar temp.AverageDriver /Karthiknew/1901.txt /output_karthiktemp2
22/06/27 15:29:41 INFO Configuration.deprecation: session.id is deprecated. Instead, use dfs.metrics.session-id
22/06/27 15:29:41 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker, sessionId=
22/06/27 15:29:41 WARN MapReduce.JobSubmitter: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
22/06/27 15:29:41 INFO input.FileInputFormat: Total input paths to process : 1
22/06/27 15:29:41 INFO mapreduce.JobSubmitter: number of splits:1
22/06/27 15:29:41 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local463464748_0001
22/06/27 15:29:41 INFO mapreduce.Job: The url to track the job: http://localhost:8888/
22/06/27 15:29:41 INFO mapreduce.Job: Running job: job_local463464748_0001
22/06/27 15:29:41 INFO mapred.LocalJobRunner: OutputCommitter set in config null
22/06/27 15:29:41 INFO mapred.LocalJobRunner: MapReduce lib: org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
22/06/27 15:29:41 INFO mapred.LocalJobRunner: Waiting for map tasks
22/06/27 15:29:41 INFO mapred.LocalJobRunner: Starting task: attempt_local463464748_0001_m_000000_0
22/06/27 15:29:41 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
22/06/27 15:29:41 INFO mapred.MapTask: Processing split: hdfs://localhost:54310/Karthiknew/1901.txt:0-888190
22/06/27 15:29:41 INFO mapred.MapTask: (EQUATOR) 0 kvt, 26214396(104857584)
22/06/27 15:29:41 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
22/06/27 15:29:41 INFO mapred.MapTask: sort limit at: 83886080
22/06/27 15:29:41 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
22/06/27 15:29:41 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
22/06/27 15:29:41 INFO mapred.MapTask: map output collector class = org.apache.hadoop.mapred.MapTask$MapOutputBuffer
22/06/27 15:29:41 INFO mapred.LocalJobRunner:
22/06/27 15:29:41 INFO mapred.MapTask: Starting flush of map output
22/06/27 15:29:41 INFO mapred.MapTask: Spilling map output
22/06/27 15:29:41 INFO mapred.MapTask: bufstart = 0; bufend = 59076; bufvoid = 104857600
22/06/27 15:29:41 INFO mapred.MapTask: kvstart = 26214396(104857584); kwend = 26188144(104752576); length = 26253/6553600
```

```
Activities Terminal Jun 27 15:51 hduser@bmsce-Precision-T1700: ~  
22/06/27 15:29:41 INFO mapred.localJobRunner: reduce > reduce  
22/06/27 15:29:41 INFO mapred.Task: Task 'attempt_local463464748_0001_r_000000_0' done.  
22/06/27 15:29:41 INFO mapred.localJobRunner: Finishing task: attempt_local463464748_0001_r_000000_0  
22/06/27 15:29:41 INFO mapred.localJobRunner: reduce task executor complete.  
22/06/27 15:29:42 INFO mapreduce.job: Job job_local463464748_0001 running in uber mode : false  
22/06/27 15:29:42 INFO mapreduce.job: map 100% reduce 100%  
22/06/27 15:29:42 INFO mapreduce.job: Job job_local463464748_0001 completed successfully  
22/06/27 15:29:42 INFO mapreduce.job: Counters: 38  
File System Counters  
FILE: Number of bytes read=176136  
FILE: Number of bytes written=746178  
FILE: Number of read operations=0  
FILE: Number of large read operations=0  
FILE: Number of write operations=0  
HDFS: Number of bytes read=1776388  
HDFS: Number of bytes written=0  
HDFS: Number of read operations=13  
HDFS: Number of large read operations=0  
HDFS: Number of write operations=4  
Map-Reduce Framework  
Map input records=6565  
Map output records=6564  
Map output bytes=59676  
Map output materialized bytes=77210  
Input split bytes=107  
Combine input records=0  
Combine output records=0  
Reduce input groups=1  
Reduce shuffle bytes=72210  
Reduce input records=6564  
Reduce output records=1  
Spilled Records=13128  
Shuffled Maps=1  
Failed Shuffles=0  
Merged Map outputs=1  
GC time elapsed (ms)=52  
CPU time spent (ms)=0  
Physical memory (bytes) snapshot=0  
Virtual memory (bytes) snapshot=0  
Total committed heap usage (bytes)=999292928  
Shuffle Errors  
BAD_ID=0  
CONNECTION=0  
IO_ERROR=0  
WRONG_LENGTH=0  
WRONG_MAP=0  
WRONG_REDUCE=0  
File Input Format Counters  
Bytes Read=888193  
File Output Format Counters  
Bytes Written=0  
hduser@bmsce-Precision-T1700:~$ hadoop fs -ls /output_karthiktemp2/part-r-00000  
-rw-r--r-- 1 hduser supergroup 0 2022-06-27 15:29 /output_karthiktemp2/part-r-00000  
hduser@bmsce-Precision-T1700:~$ hadoop fs -cat /output_karthiktemp2/part-r-00000  
1901 40  
hduser@bmsce-Precision-T1700:~$
```